



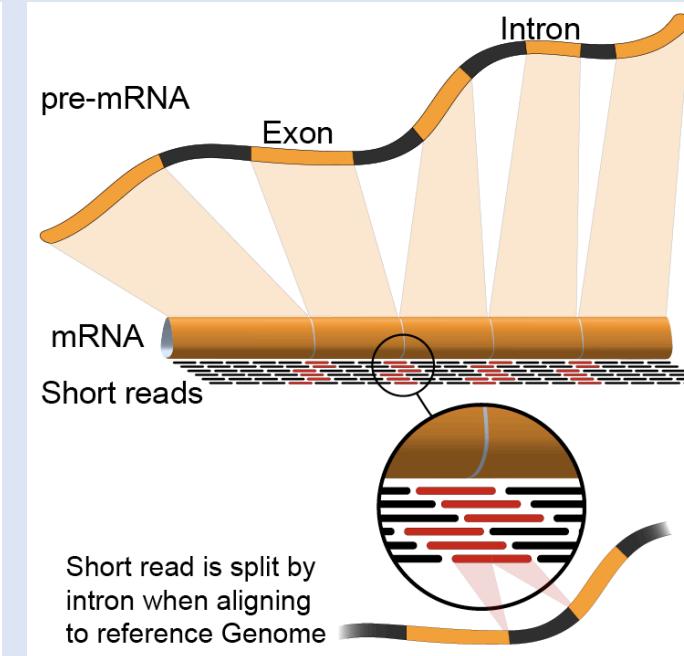
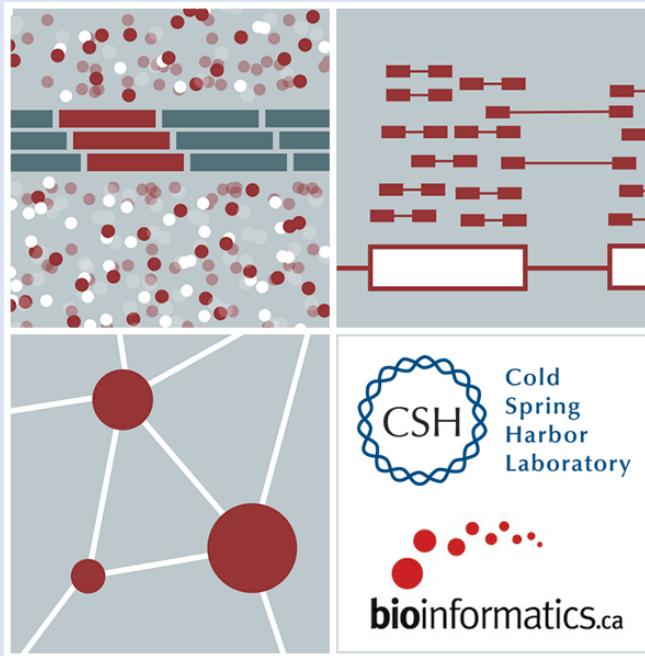
Cold
Spring
Harbor
Laboratory

RNA-Seq Module 3

Abundance Estimation and Differential Expression

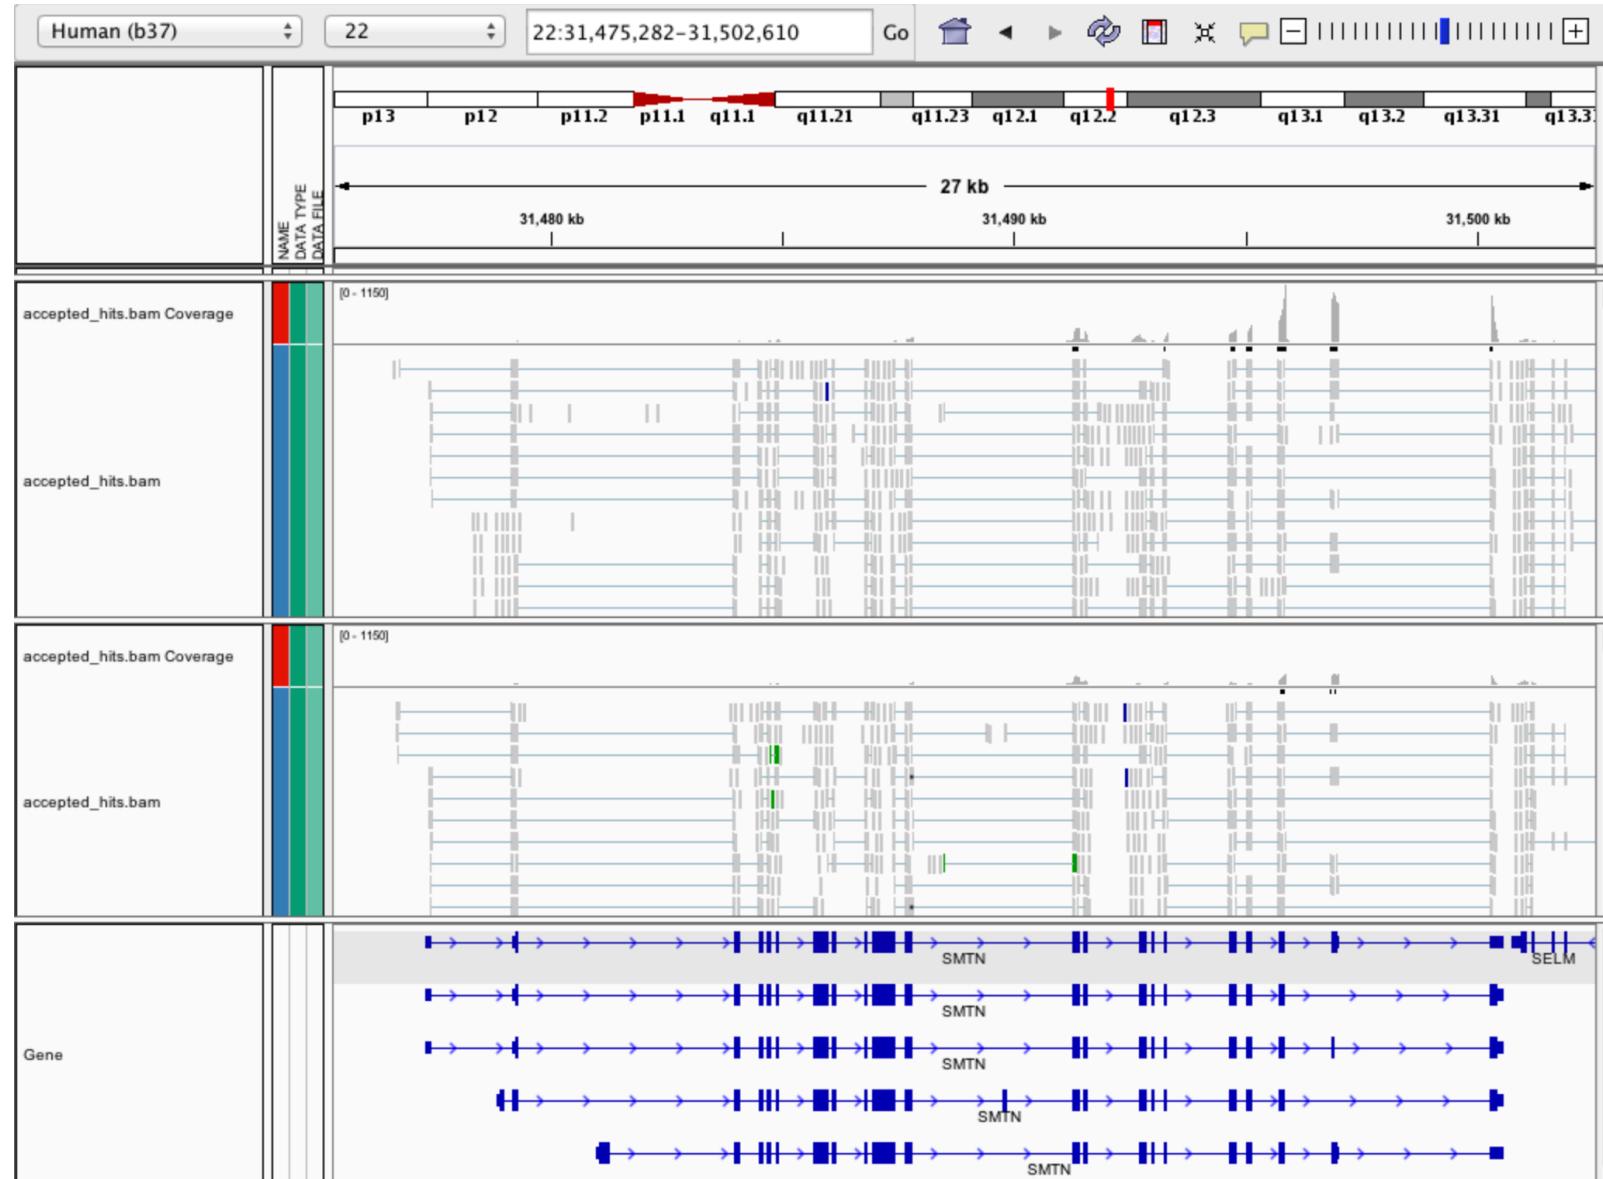
Kelsy Cotto, Arpad Danos, Harriet Dashnow, Felicia Gomez, Sharon Freshour, Obi Griffith,
Malachi Griffith, Jason Kunisaki, Chris Miller, Jonathan Preall, Aaron Quinlan

Advanced Sequencing Technologies & Bioinformatics Analysis November 11-19, 2021



Washington University in St. Louis
SCHOOL OF MEDICINE

Expression estimation for known genes and transcripts



3' bias
→

Down-regulated
↓

What is FPKM (RPKM)?

- RPKM: **Reads Per Kilobase of transcript per Million mapped reads.**
- FPKM: **Fragments Per Kilobase of transcript per Million mapped reads.**
- No essential difference - Just a terminology change to better describe paired-end reads!

What is FPKM?

- Why not just count reads in my RNAseq data? → **Fragments**
- The relative expression of a transcript is proportional to the number of cDNA fragments that originate from it. However:
 - # fragments is biased towards larger genes → **Per Kilobase of transcript**
 - # fragments is related to total library depth → **per Million mapped reads.**

What is FPKM?

- FPKM attempts to normalize for gene size and library depth
 - remember – RPKM is essentially the same!
- C = number of mappable fragments for a gene (transcript)
- N = total number of mappable fragments in the library
- L = number of base pairs in the gene (transcript)
 - $\text{FPKM} = (\text{C} / (\text{N} \times \text{L})) \times 1,000 \times 1,000,000$
 - $\text{FPKM} = (1,000,000,000 \times \text{C}) / (\text{N} \times \text{L})$
 - $\text{FPKM} = (\text{C} / (\text{N} / 1,000,000)) / (\text{L}/1000)$
- More reading:
 - <http://www.biostars.org/p/11378/>
 - <http://www.biostars.org/p/68126/>

How do FPKM and TPM differ?

- TPM: Transcript per Kilobase Million
- The difference is in the order of operations:

FPKM

- 1) Determine total fragment count, divide by 1,000,000 (per Million)
- 2) Divide each gene/transcript fragment count by #1 (Fragments Per Million)
- 3) Divide each FPM by length of each gene/transcript in kilobases (FPKM)

TPM

- 1) Divide each gene/transcript fragment count by length of the transcript in kilobases (Fragments Per Kilobase)
- 2) Sum all FPK values for the sample and divide by 1,000,000 (per Million)
- 3) Divide #1 by #2 (TPM)

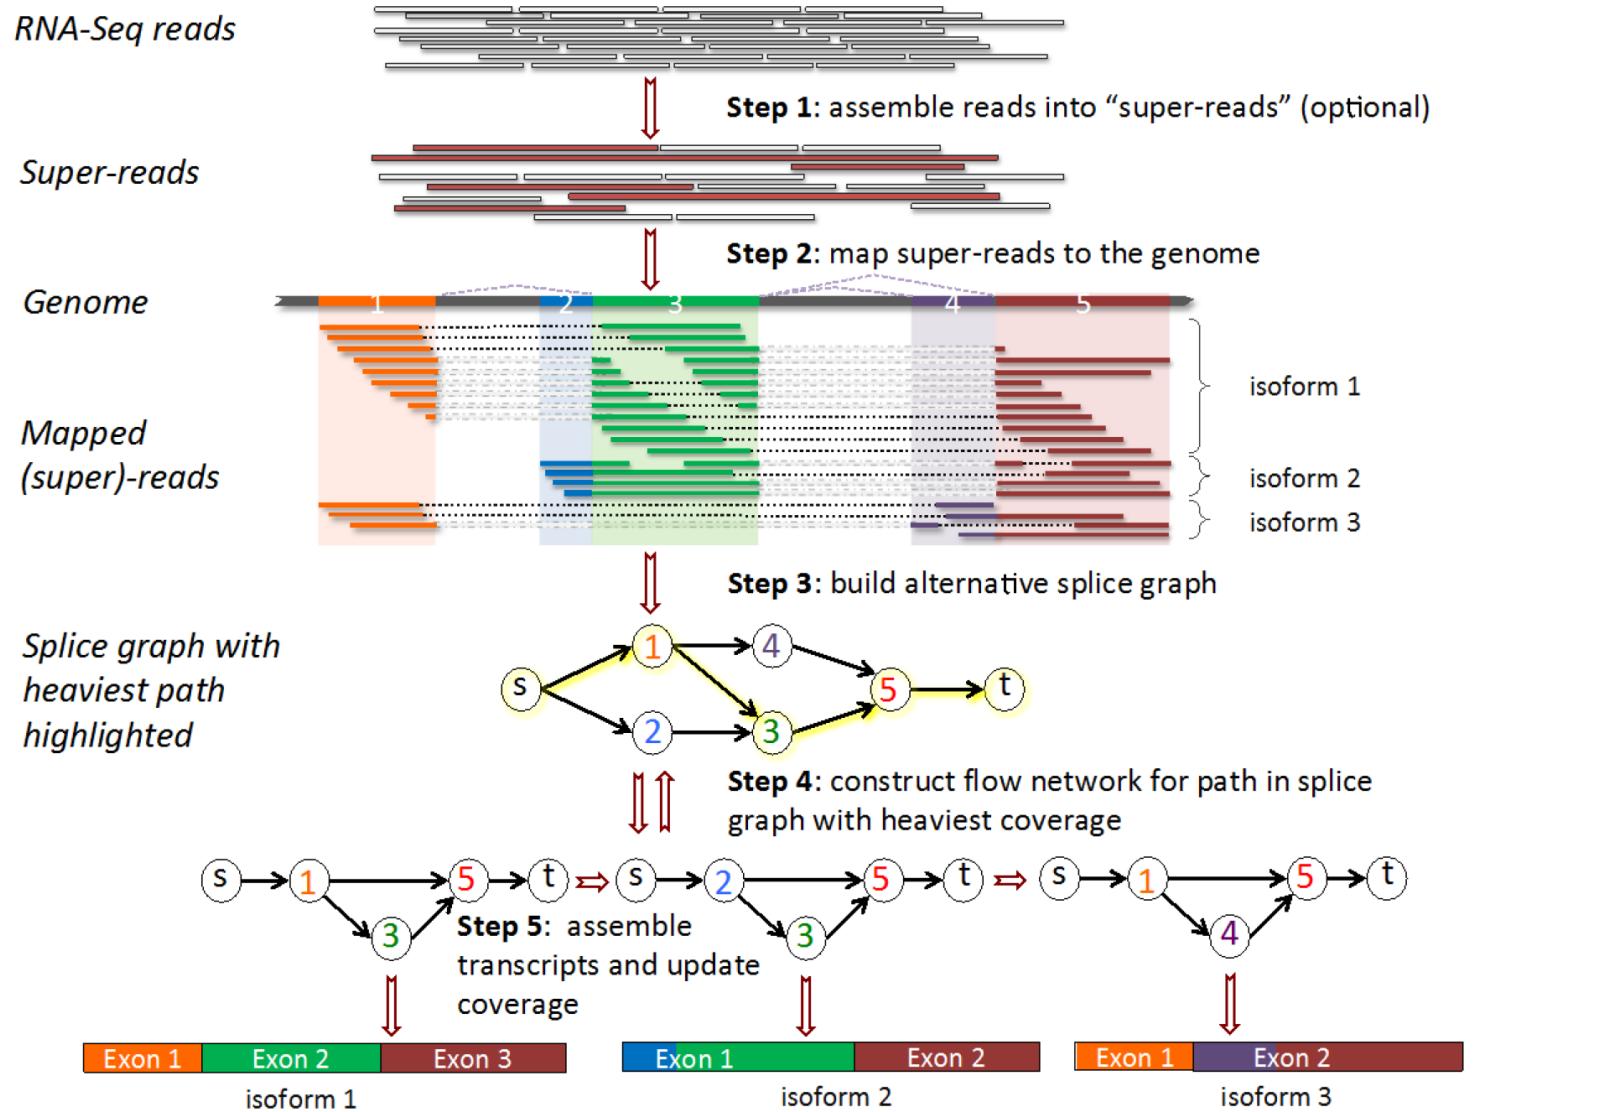
- The sum of all TPMs in each sample is the same. Easier to compare across samples!
- <http://www.rna-seqblog.com/rpkm-fpkm-and-tpm-clearly-explained/>
- <https://www.ncbi.nlm.nih.gov/pubmed/22872506>

How does StringTie work?

Map reads to the genome

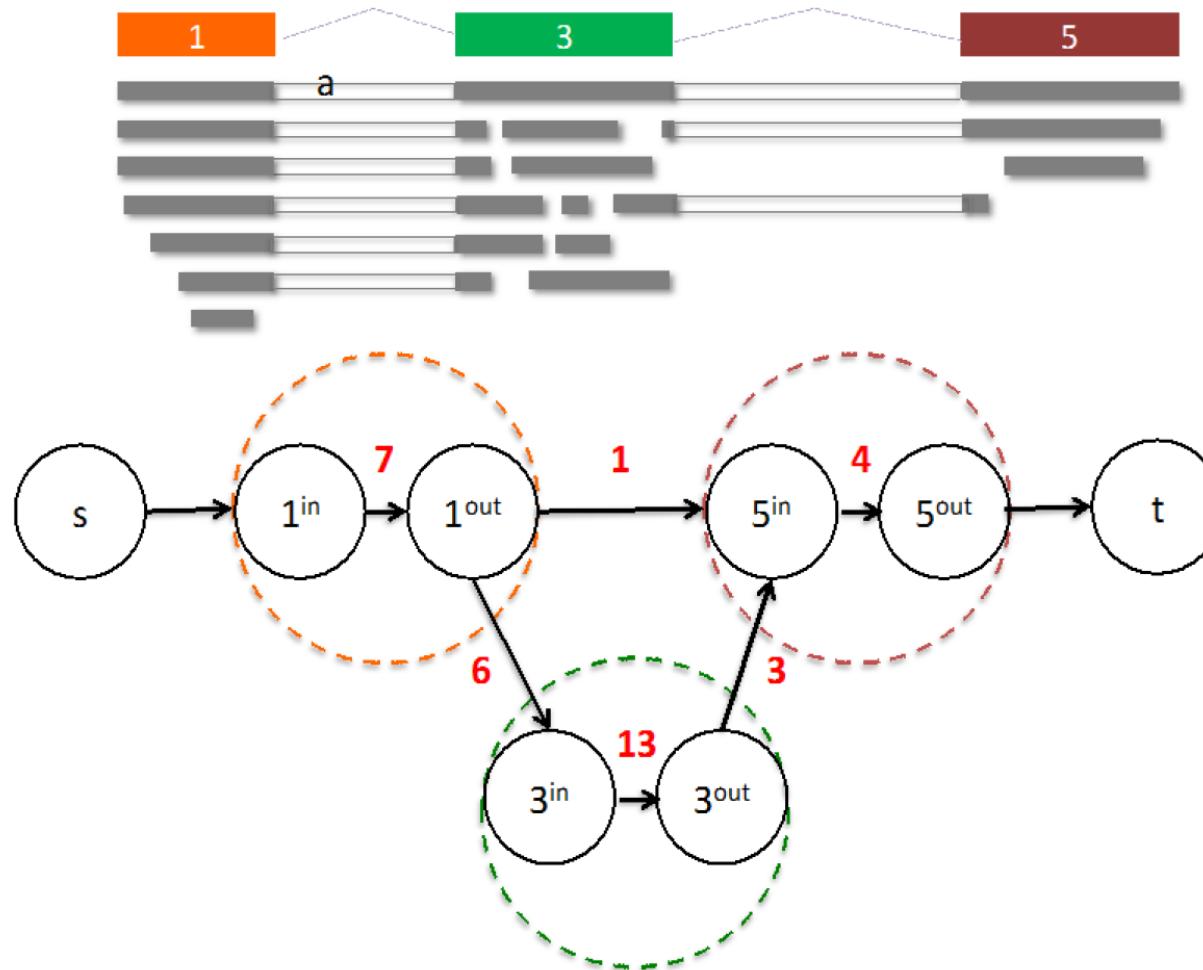
Infer isoforms:

- iteratively extract the heaviest path from a splice graph
- construct a flow network
- compute maximum flow to estimate abundance
- update the splice graph by removing reads that were assigned by the flow algorithm
- This process repeats until all reads have been assigned.



Perteau et al. Nature Biotechnology, 2015

From flow network for each transcript, maximum flow is used to assemble transcript and estimate abundance



StringTie uses basic graph theory (splice graph), custom heuristics (heaviest path), more graph theory (flow network) and optimization theory (maximum flow). See StringTie paper for definitions and math.

StringTie -merge

- Merge together all gene structures from all samples
 - Some samples may only partially represent a gene structure
- Incorporates known transcripts with assembled, potentially novel transcripts
- For de novo or reference guided mode, we will rerun StringTie with the merged transcript assembly.

Pertea et al. Nature Protocols, 2016

gffcompare

- gffcompare will compare a merged transcript GTF with known annotation, also in GTF/GFF3 format
- <http://cole-trapnell-lab.github.io/cufflinks/cuffcompare/index.html#cuffcompare-output-files>

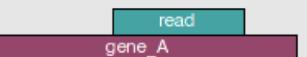
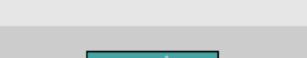
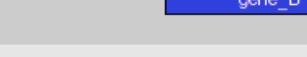
Priority	Code	Description
1	=	Complete match of intron chain
2	c	Contained
3	j	Potentially novel isoform (fragment): at least one splice junction is shared with a reference transcript
4	e	Single exon transfrag overlapping a reference exon and at least 10 bp of a reference intron, indicating a possible pre-mRNA fragment.
5	i	A transfrag falling entirely within a reference intron
6	o	Generic exonic overlap with a reference transcript
7	p	Possible polymerase run-on fragment (within 2Kbases of a reference transcript)
8	r	Repeat. Currently determined by looking at the soft-masked reference sequence and applied to transcripts where at least 50% of the bases are lower case
9	u	Unknown, intergenic transcript
10	x	Exonic overlap with reference on the opposite strand
11	s	An intron of the transfrag overlaps a reference intron on the opposite strand (likely due to read mapping errors)
12	.	(.tracking file only, indicates multiple classifications)

Alternatives to FPKM

- Raw read counts for differential expression analysis
 - Assign reads/fragments to defined genes/transcripts, get “raw counts”
 - Transcript structures could still be defined by something like Stringtie
- HTSeq (htseq-count)
 - <https://htseq.readthedocs.io/>

```
htseq-count --mode intersection-strict --stranded no --minqual 1 --type exon --idattr transcript_id accepted_hits.sam chr22.gff > transcript_read_counts_table.tsv
```
- Caveats of ‘transcript’ analysis by htseq-count:
 - Designed for genes - ambiguous reads from overlapping transcripts may not be handled!
 - <http://seqanswers.com/forums/showthread.php?t=18068>

HTSeq-count basically counts reads supporting a feature (exon, gene) by assessing overlapping coordinates

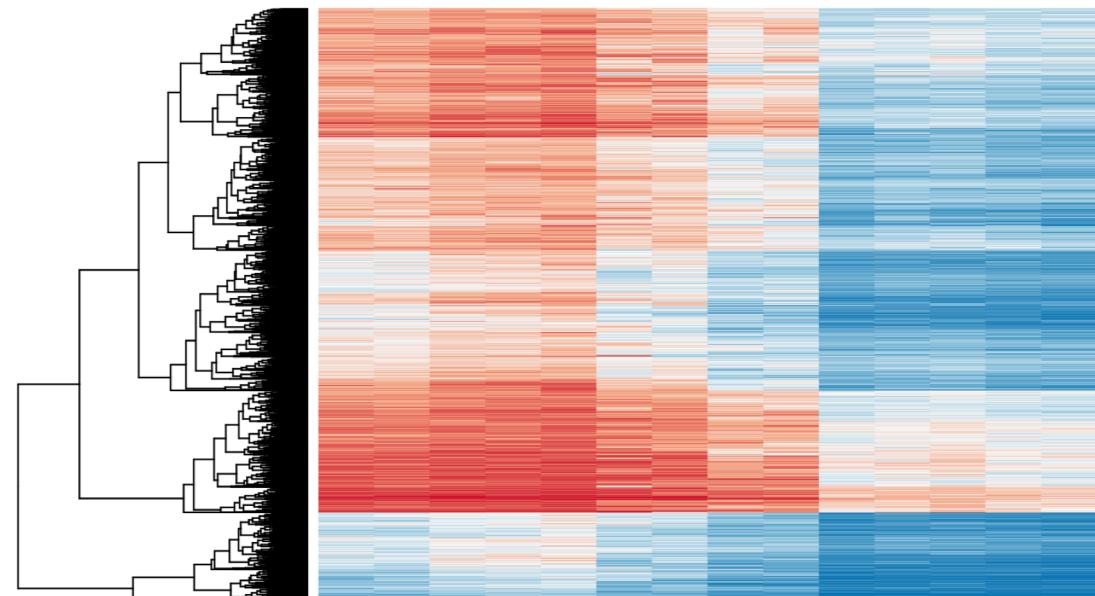
	union	intersection _strict	intersection _nonempty
 A read (teal bar) overlaps with gene_A (purple bar). The read is entirely within the gene's boundaries." data-bbox="285 255 405 295"/>	gene_A	gene_A	gene_A
 A read (teal bar) overlaps with gene_A (purple bar). The read starts after the gene's start and ends before its end." data-bbox="285 325 405 365"/>	gene_A	no_feature	gene_A
 A read (teal bar) overlaps with gene_A (purple bar) and gene_B (blue bar). The read is entirely within the gene's boundaries." data-bbox="285 405 405 445"/>	gene_A	no_feature	gene_A
 A read (teal bar) spans across gene_A (purple bar) and gene_B (blue bar). The read is entirely within the gene's boundaries." data-bbox="285 485 405 525"/>	gene_A	gene_A	gene_A
 A read (teal bar) overlaps with gene_A (purple bar) and gene_B (blue bar). The read is entirely within the gene's boundaries." data-bbox="285 565 405 605"/>	gene_A	gene_A	gene_A
 A read (teal bar) overlaps with gene_A (purple bar) and gene_B (blue bar). The read is entirely within the gene's boundaries." data-bbox="285 645 405 685"/>	ambiguous	gene_A	gene_A
 A read (teal bar) overlaps with gene_A (purple bar) and gene_B (blue bar). The read is entirely within the gene's boundaries." data-bbox="285 725 405 765"/>	ambiguous	ambiguous	ambiguous

Note, if gene_A and gene_B on opposite strands, sequence data is stranded, and correct HTSeq parameter set then this read may not be ambiguous

Whether a read is counted depends on the nature of overlap and "mode" selected

Differential Expression

- Tying gene expression back to genotype/phenotype
- What genes/transcripts are being expressed at higher/lower levels in different groups of samples?
 - Are these differences 'significant', accounting for variance/noise?
- Examples (used in course):
 - UHR cells vs HBR brain
 - Tumor vs Normal tissue
 - Wild-type vs gene KO cells



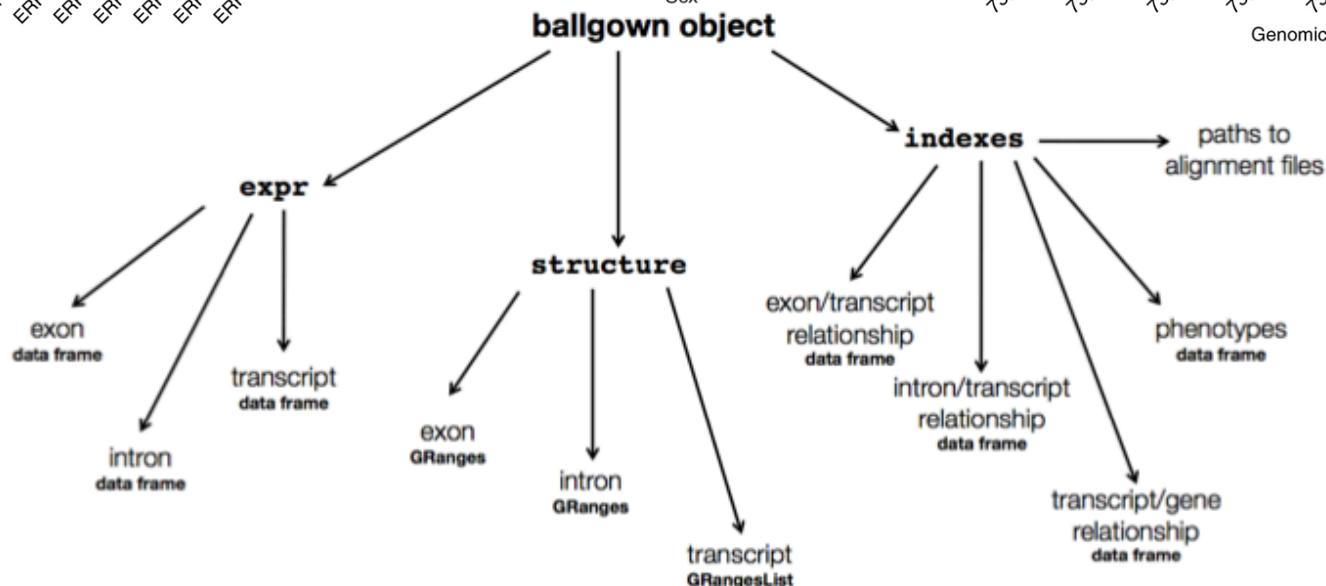
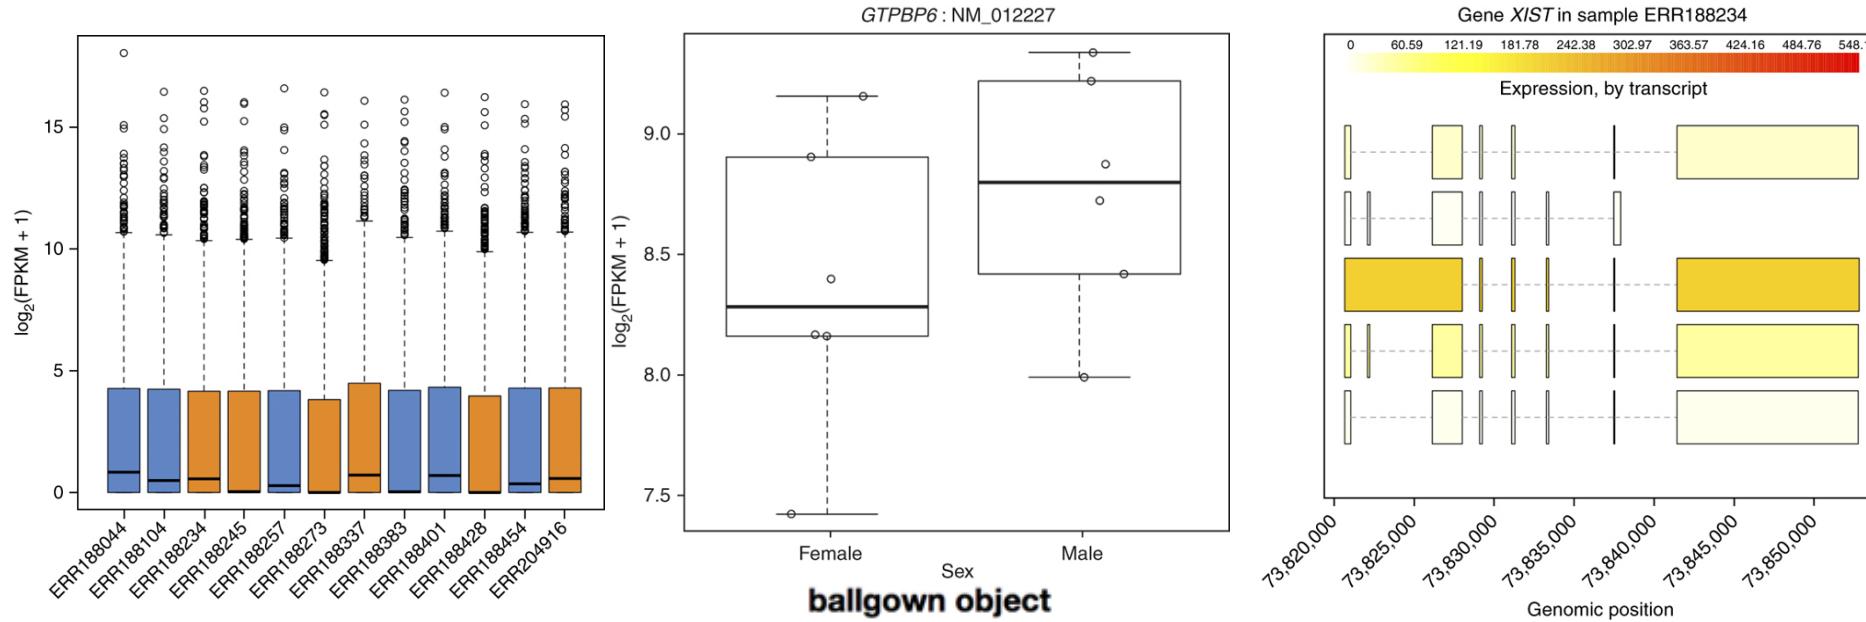
Differential Expression with Ballgown

Parametric F-test comparing nested linear models

- Two models are fit to each feature, using expression as the outcome
 - one including the covariate of interest (e.g., case/control status or time) and one not including that covariate.
- An F statistic and p-value are calculated using the fits of the two models.
 - A significant p-value means the model including the covariate of interest fits significantly better than the model without that covariate, indicating differential expression.
- We adjust for multiple testing by reporting q-values:
 - $q < 0.05$ the false discovery rate should be controlled at $\sim 5\%$.

[Frazee et al. \(2014\)](#)

Ballgown for Visualization with R



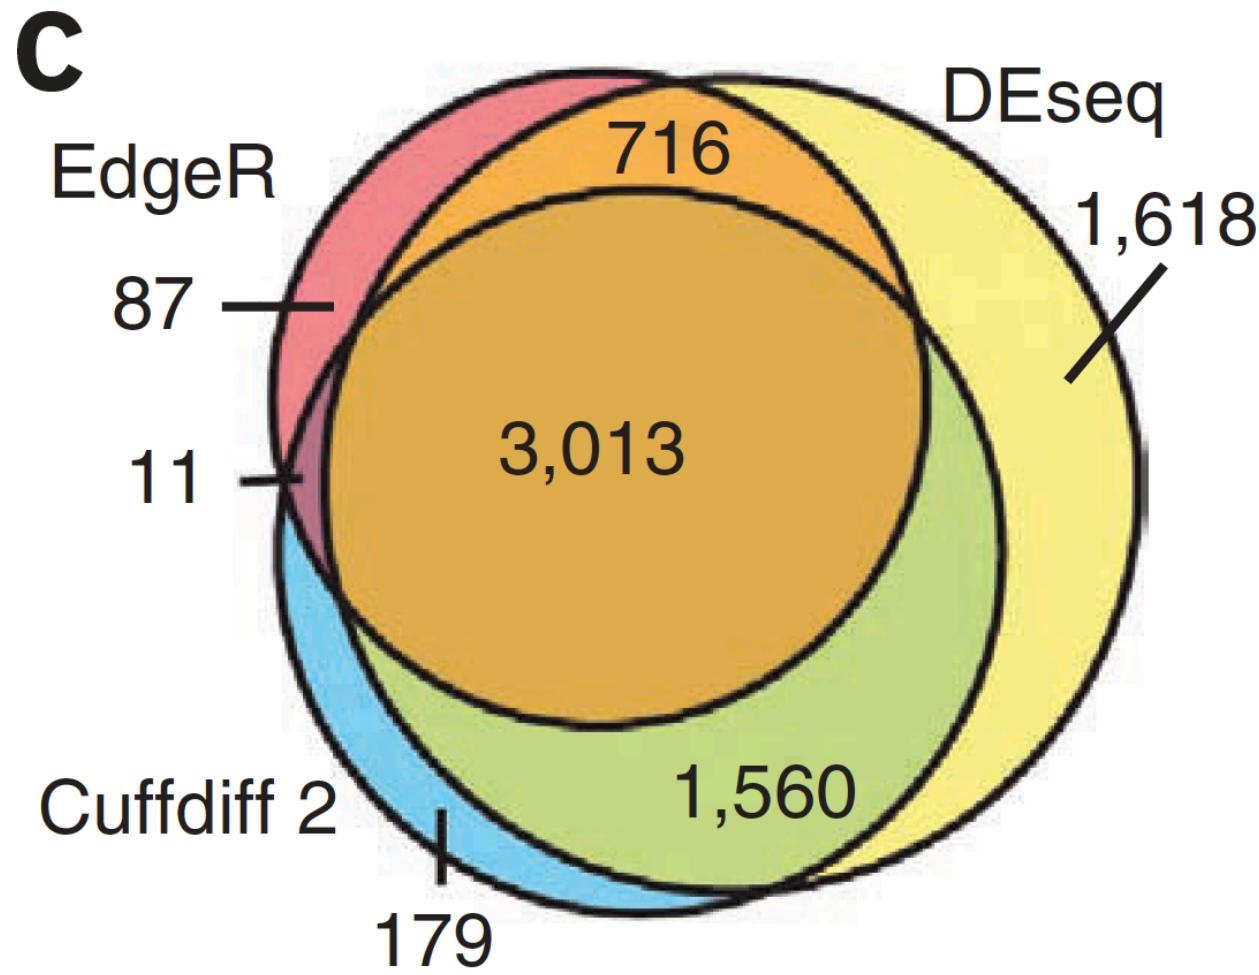
Alternative differential expression methods

- Raw count approaches
 - DESeq2 - <http://www-huber.embl.de/users/anders/DESeq/>
 - edgeR - <http://www.bioconductor.org/packages/release/bioc/html/edgeR.html>
 - Others...

'FPKM/TPM' expression estimates vs. 'raw' counts

- Which should I use?
 - Long running debate, but the general consensus:
- FPKM/TPM
 - When you want to leverage benefits of tuxedo suite
 - Isoform deconvolution
 - Good for visualization (e.g., heatmaps)
 - Calculating fold changes, etc.
- Counts
 - More robust statistical methods for differential expression
 - Accommodates more sophisticated experimental designs with appropriate statistical tests

Multiple approaches advisable



Lessons learned from microarray days

- Hansen et al. “Sequencing Technology Does Not Eliminate Biological Variability.” *Nature Biotechnology* 29, no. 7 (2011): 572–573.
- Power analysis for RNA-seq experiments
 - <http://scotty.genetics.utah.edu/>
- RNA-seq need for biological replicates
 - <http://www.biostars.org/p/1161/>
- RNA-seq study design
 - <http://www.biostars.org/p/68885/>

Multiple testing correction

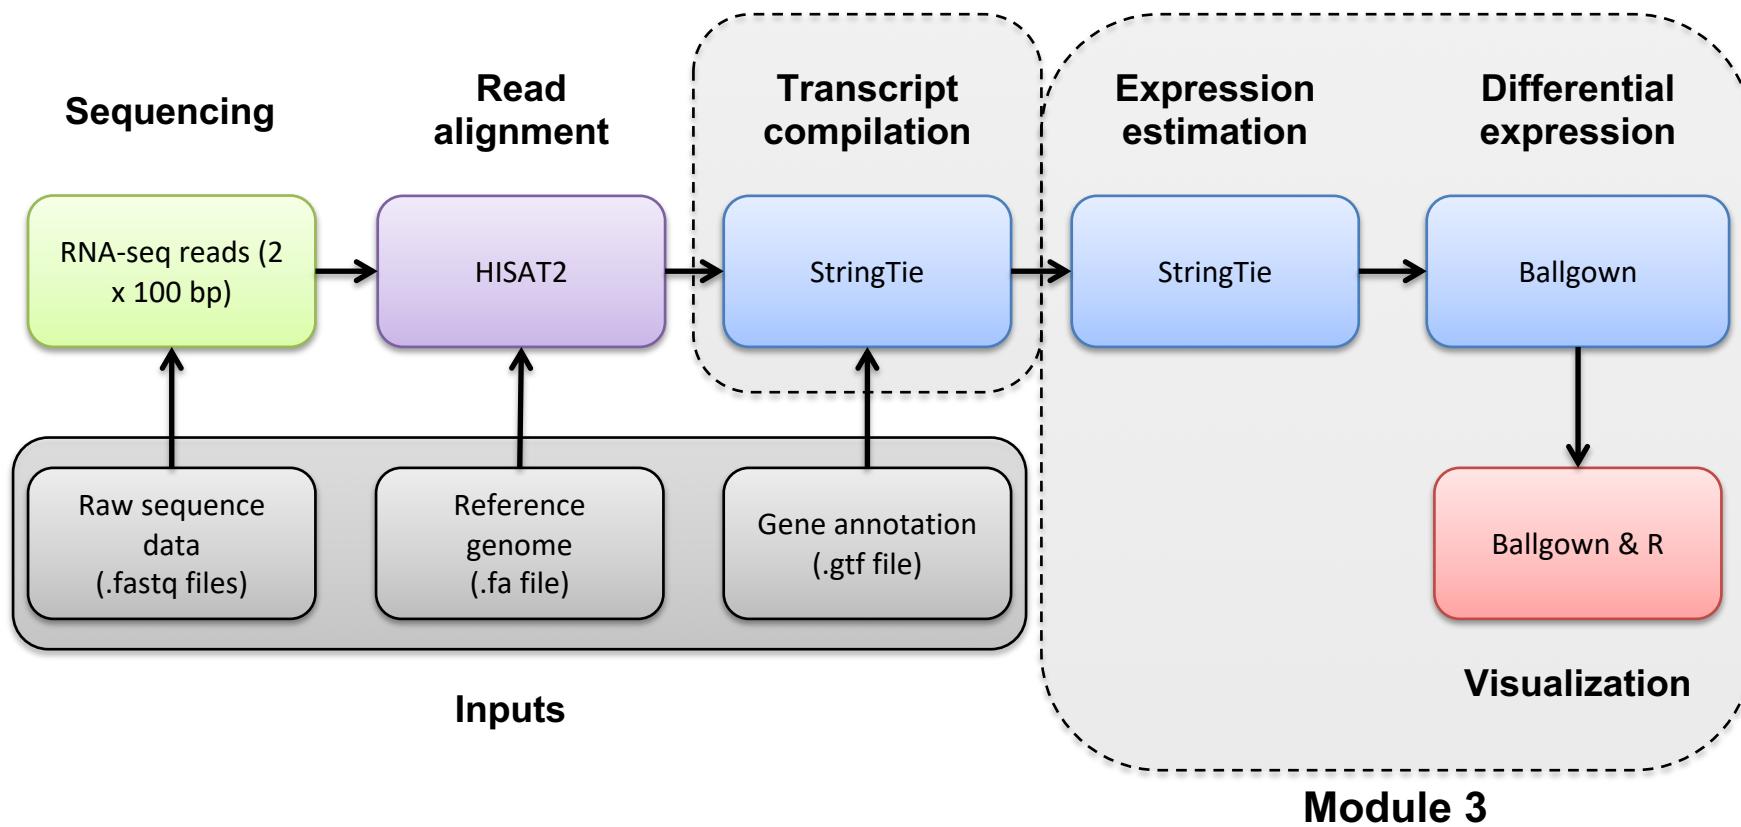
- As more attributes are compared, differences due solely to chance become more likely!
- Well known from array studies
 - 10,000s genes/transcripts
 - 100,000s exons
- With RNA-seq, more of a problem than ever
 - All the complexity of the transcriptome gives huge numbers of potential features
 - Genes, transcripts, exons, junctions, retained introns, microRNAs, lncRNAs, etc
- Bioconductor multtest
 - <http://www.bioconductor.org/packages/release/bioc/html/multtest.html>

Downstream interpretation of expression analysis

- Topic for an entire course
- Expression estimates and differential expression lists from StringTie, Ballgown or other alternatives can be fed into many analysis pipelines
- See supplemental R tutorial for how to format expression data and start manipulating in R
- Clustering/Heatmaps
 - Provided by Ballgown
 - For more customized analysis various R packages exist:
 - hclust, heatmap.2, plotrix, ggplot2, etc.
- Classification
 - For RNA-seq data we still rarely have sufficient sample size and clinical details but this is changing
 - Weka is a good learning tool
 - RandomForests R package (biostar tutorial being developed)
- Pathway analysis
 - GSEA, IPA, Cytoscape, many R/BioConductor packages:
<http://www.bioconductor.org/help/search/index.html?q=pathway>

https://genviz.org/module-04-expression/0004/01/01/Expression_Profiling_and_Visualization/

HISAT2/StringTie/Ballgown RNA-seq Pipeline



We are on a Coffee Break &
Networking Session