

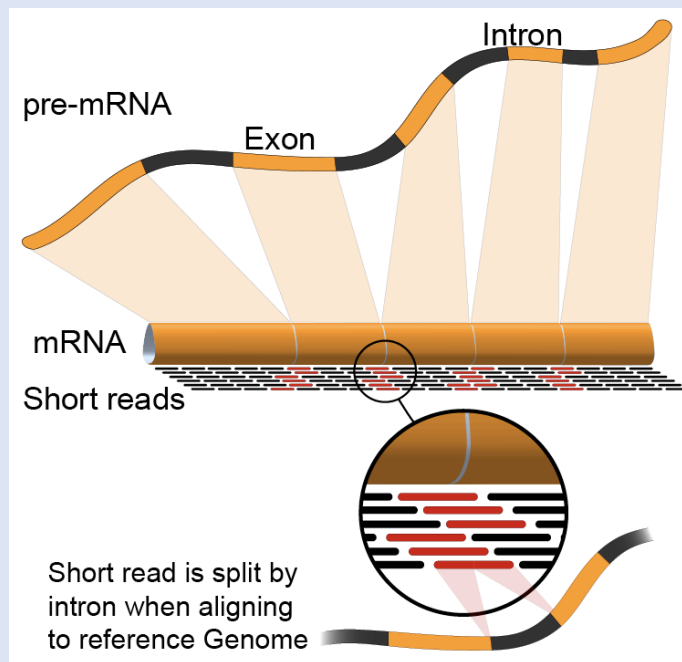


Cold  
Spring  
Harbor  
Laboratory

# RNA-Seq Module 2: SAM/BAM/BED file formats

Arpad Danos, Felicia Gomez, Obi Griffith, Malachi Griffith,  
My Hoang, Mariam Khanfar, Chris Miller, Kartik Singhal

Advanced Sequencing Technologies & Bioinformatics Analysis November 5-19, 2023



Washington University in St. Louis  
SCHOOL OF MEDICINE

# Introduction to the SAM/BAM format

- The specification
  - <http://samtools.sourceforge.net/SAM1.pdf>
- SAM is uncompressed text data
- BAM is a compressed version of SAM
  - lossless BGZF format
- BAM files are usually ‘indexed’
  - A ‘.bai’ file will be found beside the ‘.bam’ file
  - Indexing provides fast retrieval of alignments overlapping a specified region without going through all alignments.
  - BAM must be sorted by the reference ID and then the leftmost coordinate before indexing

# Example of SAM/BAM file format

## Example SAM/BAM/CRAM header section (abbreviated)

```
mqriffit@linux270 ~$ samtools view -H /gscmnt/gc13001/info/model_data/2891632684/build136494552/alignments/136080019.bam | grep -P "SN:[2\|HD\|RG\|PG"
```

```
@HD      VN:1.4    SO:coordinate  
@SQ      SN:12   LN:51304566     UR:ftp://ftp.ncbi.nih.gov/genbank/genomes/Eukaryotes/vertebrates_mammals/Homo_sapiens/GRCh37/special_requests/GRCh37-lite.fa.gz AS:GRCh37-lite M5:a718acaa6135fdca8357d5bfe9  
4211dd   SP:Homo sapiens  
@RG      ID:2888721359   PL:illumina   PI:D1B4ACXX.3 LB:H_KA-452198-0817007-cDNA-3-lib1   PI:365 DS:paired end DT:2012-10-03T19:00:00-0500 SM:H_KA-452198-0817007 ON:WUGSC  
@PG      ID:2888721359   VN:2.0.8     CL:tophat --library-type fr-secondstrand --bowtie-version=2.1.0  
@PG      ID:MarkDuplicates PN:MarkDuplicates PP:2888721359 VN:1.85(exported) CL:net.sf.picard.sam.MarkDuplicates INPUT=[/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwaker-15434-136080019/scrch-Ilg6v/H_KA-452198-0817007-cDNA-3-lib1-2888360300.bam] OUTPUT=/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwaker-15434-136080019/scrch-Ilg6v/H_KA-452198-0817007-cDNA-3-lib1-2888360300-post.dup.bam METRICS_FILE=/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwaker-15434-136080019/staging-ljuJS/H_KA-452198-0817007-cDNA-3-lib1-2888360300.metrics REMOVE_DUPLICATES=false ASSUME_SORTED=true MAX_FILE_HANDLES_FOR_READ_ENDS_MAP=9500 TMP_DIR=[/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwaker-15434-136080019/scrch-Ilg6v] VALIDATION_STRINGENCY=SILENT MAX_RECORDS_IN_RAM=50000 PROGRAM_RECORD_ID=MarkDuplicates PROGRAM_GROUP_NAME=Mark Duplicates MAX_SEQUENCES_FOR_DISK_READ_ENDS_MAP=50000 SORTING_COLLECTION_SIZE_RATIO=0.25 READ_NAME_REGEX=[a-zA-Z0-9+]+:(0-9)+:([0-9+]):([0-9+]):([0-9+]).* OPTICAL_DUPLICATE_PIXEL_DISTANCE=100 VERBOSITY=INFO QUIET=false COMPRESSION_LEVEL=5 CREATE_INDEX=false CREATE_MD5_FILE=false
```

```
mqriffit@linux270 ~$
```

## Example SAM/BAM/CRAM alignment section (only 10 alignments shown)

[illegible]

# SAM/BAM header section

- Used to describe source of data, reference sequence, method of alignment, etc.
- Each section begins with character '@' followed by a two-letter record type code. These are followed by two-letter tags and values:
  - @HD The header line
    - VN: format version
    - SO: Sorting order of alignments
  - @SQ Reference sequence dictionary
    - SN: reference sequence name
    - LN: reference sequence length
    - SP: species
  - @RG Read group
    - ID: read group identifier
    - CN: name of sequencing center
    - SM: sample name
  - @PG Program
    - PN: program name
    - VN: program version

# A BAM file is divided in header and alignment sections

## Example SAM/BAM header section (abbreviated)

```
mgriffit@linus270 ~$ samtools view -H /gscmnt/gc13001/info/model_data/2891632684/build136494552/alignments/136080019.bam | grep -P "SN\.:22|HD|RG|PG"
@HD      VN:1.4  SO:coordinate
@SQ      SN:22  LN:51304566      UR:ftp://ftp.ncbi.nih.gov/genbank/genomes/Eukaryotes/vertebrates_mammals/Homo_sapiens/GRCh37/special_requests/GRCh37-lite.fa.gz AS:GRCh37-lite M5:a718acaa6135fdca8357d5bfe9
4211dd SP:Homo sapiens
@RG      ID:2888721359  PL:illumina  PU:D1BA4ACXX.3  LB:H_KA-452198-0817007-cDNA-3-lib1  PI:365  DS:paired end  DT:2012-10-03T19:00:00-0500  SM:H_KA-452198-0817007  CN:WUGSC
@PG      ID:2888721359  VN:2.0.8      CL:tophat --library-type fr-secondstrand --bowtie-version=2.1.0
@PG      ID:MarkDuplications  PN:MarkDuplications  PP:2888721359  VN:1.85(exported)  CL:net.sf.picard.sam.MarkDuplications INPUT=[/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwalker-15434-136080019/scratch-ILg6Y/H_KA-452198-0817007-cDNA-3-lib1-2888360300.bam] OUTPUT=/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwalker-15434-136080019/scratch-ILg6Y/H_KA-452198-0817007-cDNA-3-lib1-2888360300-post_dup.bam METRICS_FILE=/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwalker-15434-136080019/staging-1iuJS/H_KA-452198-0817007-cDNA-3-lib1-2888360300.metrics REMOVE_DUPLICATES=false ASSUME_SORTED=true MAX_FILE_HANDLES_FOR_READ_ENDS_MAP=9500 TMP_DIR=[/gscmnt/gc13001/info/build_merged_alignments/merged-alignment-blade10-2-5.gsc.wustl.edu-jwalker-15434-136080019/scratch-ILg6Y] VALIDATION_STRINGENCY=SILENT MAX_RECORDS_IN_RAM=500000 PROGRAM_RECORD_ID=MarkDuplications PROGRAM_GROUP_NAME=MarkDuplications MAX_SEQUENCES_FOR_DISK_READ_ENDS_MAP=50000 SORTING_COLLECTION_SIZE_RATIO=0.25 READ_NAME_REGEX=[a-zA-Z0-9]+:[0-9]:([0-9]+):([0-9]+):([0-9]+):([0-9]+).* OPTICAL_DUPLICATE_PIXEL_DISTANCE=100 VERBOSITY=INFO QUIET=false COMPRESSION_LEVEL=5 CREATE_INDEX=false CREATE_MD5_FILE=false
mgriffit@linus270 ~$
```

Version (VN) and sort order (SO) - Important!

Reference sequence (SQ) and sequence length (LN)

```
@HD      VN:1.3  SO:coordinate
@SQ      SN:20  LN:63025520
@RG      ID:HG00096  SM:HG00096
@PG      ID:HG00096  PN:bwa  CL:/Users/AlistairNWard/Work/gkno/gkno_launcher/tools/bwa/bwa mem -t 4
```

Read group (RG) and sample (SM)

Programs (PG) that have been run on the data



A BAM file is divided in header and alignment sections

# Example SAM/BAM alignment section (only 10 alignments shown)

[illegible]

# SAM/BAM alignment section

| Col | Field | Type   | Regex/Range                              | Brief description                     |
|-----|-------|--------|--|---------------------------------------|
| 1   | QNAME | String | [!-?A-~]{1,255}                          | Query template NAME                   |
| 2   | FLAG  | Int    | [0,2 <sup>16</sup> -1]                   | bitwise FLAG                          |
| 3   | RNAME | String | \*  [!-( )+-<>-~] [!-~]*                 | Reference sequence NAME               |
| 4   | POS   | Int    | [0,2 <sup>29</sup> -1]                   | 1-based leftmost mapping POSition     |
| 5   | MAPQ  | Int    | [0,2 <sup>8</sup> -1]                    | MAPping Quality                       |
| 6   | CIGAR | String | \*  ([0-9]+[MIDNSHPX=])+                 | CIGAR string                          |
| 7   | RNEXT | String | \* =  [!-( )+-<>-~] [!-~]*               | Ref. name of the mate/next segment    |
| 8   | PNEXT | Int    | [0,2 <sup>29</sup> -1]                   | Position of the mate/next segment     |
| 9   | TLEN  | Int    | [-2 <sup>29</sup> +1,2 <sup>29</sup> -1] | observed Template LENgth              |
| 10  | SEQ   | String | \*  [A-Za-z=.]+                          | segment SEQUENCE                      |
| 11  | QUAL  | String | [!-~]+                                   | ASCII of Phred-scaled base QUALity+33 |

## Example values

|    |       |      |   |
|----|-------|------|---|
| 1  | QNAME | e.g. | HWI-ST495_129147882:1:2302:10269:12362  |
| 2  | FLAG  | e.g. | 99  |
| 3  | RNAME | e.g. | 1   |
| 4  | POS   | e.g. | 11623   |
| 5  | MAPQ  | e.g. | 3   |
| 6  | CIGAR | e.g. | 100M  |
| 7  | RNEXT | e.g. | =   |
| 8  | PNEXT | e.g. | 11740   |
| 9  | TLEN  | e.g. | 217   |
| 10 | SEQ   | e.g. | CCTGTTTCTCCACAAAGTGTTTACTTTTGGATTTTTGCCAGTCTAACAGGTGAAGCCCTGGAGATTCTTATTAGTGATTTGGGCTGGGGCCTGGCCATGT    |
| 11 | QUAL  | e.g. | CCCFFFFFFHHHHHJJJIJFIJJJJJJJJJJJHIJJJJJJJIJJJJJGGHIJHIJJJJJJJJJGHGGIJJJJJJIJEEHHHHFFFFCDCCCCDDDDDB@ACDD |

# SAM Format – Information Fields

| Col | Field | Type   | Regex/Range                              | Brief description                     |
|-----|-------|--------|--|---------------------------------------|
| 1   | QNAME | String | [!-?A-~]{1,255}                          | Query template NAME                   |
| 2   | FLAG  | Int    | [0,2 <sup>16</sup> -1]                   | bitwise FLAG                          |
| 3   | RNAME | String | \*  [!-()+-<>-~] [!-~]*                  | Reference sequence NAME               |
| 4   | POS   | Int    | [0,2 <sup>31</sup> -1]                   | 1-based leftmost mapping POSition     |
| 5   | MAPQ  | Int    | [0,2 <sup>8</sup> -1]                    | MAPping Quality                       |
| 6   | CIGAR | String | \*  ([0-9]+[MIDNSHPX=])+                 | CIGAR string                          |
| 7   | RNEXT | String | \* =  [!-()+-<>-~] [!-~]*                | Ref. name of the mate/next read       |
| 8   | PNEXT | Int    | [0,2 <sup>31</sup> -1]                   | Position of the mate/next read        |
| 9   | TLEN  | Int    | [-2 <sup>31</sup> +1,2 <sup>31</sup> -1] | observed Template LENgth              |
| 10  | SEQ   | String | \*  [A-Za-z=.]+                          | segment SEQuence                      |
| 11  | QUAL  | String | [!-~]+                                   | ASCII of Phred-scaled base QUALity+33 |

|                    |     |    |        |    |      |   |        |     |  |
|--------------------|-----|----|--------|----|------|---|--------|-----|--|
| 1                  | 2   | 3  | 4      | 5  | 6    | 7 | 8      | 9   | 10                                     |
| SRR062634.14576120 | 163 | 20 | 899919 | 60 | 100M | = | 900037 | 218 | TTCCCCAGTAGCTGGGATTACAGGCATACGCCACCATC |
|                    | ?   |    |        |    | ?    |   |        |     |  |



# SAM/BAM flags explained

- 12 bitwise flags describing the alignment
- Stored as a binary string of length 12 instead of 12 columns of data
- Value of '1' indicates the flag is set. e.g. 001000000000
- All combinations can be represented as a number from 0 to 4095 (i.e.  $2^{12}-1$ ). This number is used in the BAM/SAM file.
- You can specify 'required' or 'filter' flags in samtools view using the '-f' and '-F' options respectively

| Bit  |       | Description  |
|------|-------|--|
| 1    | 0x1   | template having multiple segments in sequencing                    |
| 2    | 0x2   | each segment properly aligned according to the aligner             |
| 4    | 0x4   | segment unmapped   |
| 8    | 0x8   | next segment in the template unmapped                              |
| 16   | 0x10  | SEQ being reverse complemented                                     |
| 32   | 0x20  | SEQ of the next segment in the template being reverse complemented |
| 64   | 0x40  | the first segment in the template                                  |
| 128  | 0x80  | the last segment in the template                                   |
| 256  | 0x100 | secondary alignment  |
| 512  | 0x200 | not passing filters, such as platform/vendor quality controls      |
| 1024 | 0x400 | PCR or optical duplicate   |
| 2048 | 0x800 | supplementary alignment  |

Note that to maximize confusion, each bit is described in the SAM specification using its hexadecimal representation (i.e., '0x10' = 16 and '0x40' = 64).

<http://broadinstitute.github.io/picard/explain-flags.html>

# SAM Format – Information Fields

| Col | Field | Type   | Regex/Range                              | Brief description                     |
|-----|-------|--------|--|---------------------------------------|
| 1   | QNAME | String | [!-?A-~]{1,255}                          | Query template NAME                   |
| 2   | FLAG  | Int    | [0,2 <sup>16</sup> -1]                   | bitwise FLAG                          |
| 3   | RNAME | String | \*  [!-( )+-<>-~] [!-~]*                 | Reference sequence NAME               |
| 4   | POS   | Int    | [0,2 <sup>31</sup> -1]                   | 1-based leftmost mapping POSition     |
| 5   | MAPQ  | Int    | [0,2 <sup>8</sup> -1]                    | MAPping Quality                       |
| 6   | CIGAR | String | \*  ([0-9]+[MIDNSHPX=])+                 | CIGAR string                          |
| 7   | RNEXT | String | \* =  [!-( )+-<>-~] [!-~]*               | Ref. name of the mate/next read       |
| 8   | PNEXT | Int    | [0,2 <sup>31</sup> -1]                   | Position of the mate/next read        |
| 9   | TLEN  | Int    | [-2 <sup>31</sup> +1,2 <sup>31</sup> -1] | observed Template LENgth              |
| 10  | SEQ   | String | \*  [A-Za-z=.]+                          | segment SEQUENCE                      |
| 11  | QUAL  | String | [!-~]+                                   | ASCII of Phred-scaled base QUALity+33 |

|                    |     |    |        |    |      |   |        |     |  |
|--------------------|-----|----|--------|----|------|---|--------|-----|--|
| 1                  | 2   | 3  | 4      | 5  | 6    | 7 | 8      | 9   | 10                                     |
| SRR062634.14576120 | 163 | 20 | 899919 | 60 | 100M | = | 900037 | 218 | TTCCCCAGTAGCTGGGATTACAGGCATACGCCACCATC |
|                    | ?   |    |        |    | ?    |   |        |     |  |

# CIGAR strings explained

- The 'CIGAR' (Compact Idiosyncratic Gapped Alignment Report)
- The CIGAR string is a sequence of base lengths and associated 'operations' indicating which bases align to the reference (either a match or mismatch), are deleted, are inserted, represent introns, etc.

| Op | BAM | Description   |
|----|-----|---|
| M  | 0   | alignment match (can be a sequence match or mismatch) |
| I  | 1   | insertion to the reference                            |
| D  | 2   | deletion from the reference                           |
| N  | 3   | skipped region from the reference                     |
| S  | 4   | soft clipping (clipped sequences present in SEQ)      |
| H  | 5   | hard clipping (clipped sequences NOT present in SEQ)  |
| P  | 6   | padding (silent deletion from padded reference)       |
| =  | 7   | sequence match  |
| X  | 8   | sequence mismatch                                     |

- e.g. 81M859N19M

- A 100 bp read consists of: 81 bases of alignment to reference, 859 bases skipped (an intron), 19 bases of alignment

# CRAM files

- CRAM is an ultra-compressed version of a BAM file
  - Usually between 30-60% smaller than the corresponding BAM
- Stores “diffs” from the reference genome
  - requires the matching reference genome to restore original data!
- Base quality binning may be used as well
- Some tools still require conversion back to bam

| Quality Score Bins | Example of Empirically Mapped Quality Scores* |
|--------------------|---|
| N (no call)        | N (no call)                                   |
| 2-9                | 6   |
| 10-19              | 15  |
| 20-24              | 22  |
| 25-29              | 27  |
| 30-34              | 33  |
| 35-39              | 37  |
| ≥ 40               | 40  |

By replacing the quality scores between 19 and 25 with a new score of 22, data storage space is conserved.

\*The mapped quality score of each bin (except “N”) is subject to change depending on individual Q-tables.



# Introduction to the BED format

- When working with BAM files, it is very common to want to examine a focused subset of the reference genome
  - e.g. the exons of a gene
- These subsets are commonly specified in 'BED' files
  - <https://genome.ucsc.edu/FAQ/FAQformat.html#format1>
- Many BAM manipulation tools accept regions of interest in BED format
- Basic BED format (tab separated):
  - Chromosome name, start position, end position (BED3)
  - Coordinates in BED format are 0 based

# Introduction to the BED format

- There are several flavors of BED format: BED3, BED4, BED6, BED8, etc
- First 3 fields always required: chr, start, stop
- Followed by up to 9 additional optional fields: name, score, strand, thickStart, thickEnd, itemRGB, blockCount, blockSizes, blockStarts

|      |           |           |      |   |   |
|------|-----------|-----------|------|---|---|
| chr7 | 127471196 | 127472363 | Pos1 | 0 | + |
| chr7 | 127472363 | 127473530 | Pos2 | 0 | + |
| chr7 | 127473530 | 127474697 | Pos3 | 0 | + |
| chr7 | 127474697 | 127475864 | Pos4 | 0 | + |
| chr7 | 127475864 | 127477031 | Neg1 | 0 | - |
| chr7 | 127477031 | 127478198 | Neg2 | 0 | - |
| chr7 | 127478198 | 127479365 | Neg3 | 0 | - |
| chr7 | 127479365 | 127480532 | Pos5 | 0 | + |
| chr7 | 127480532 | 127481699 | Neg4 | 0 | - |

# Manipulation of SAM/BAM and BED files

- Several tools are used ubiquitously in sequence analysis to manipulate these files
- SAM/BAM files
  - samtools
  - bamtools
  - Picard
- BED files
  - bedtools
  - bedops




# Common sources of confusion

- Genomic coordinate systems
- Genome builds
- Variant representation



# Genomic coordinates – 1 vs 0 based



|                                      | 1-based      | 0-based      |
|--------------------------------------|--------------|--------------|
| Indicate a single nucleotide         | chr1:4-4 G   | chr1:3-4 G   |
| Indicate a range of nucleotides      | chr1:2-4 ACG | chr1:1-4 ACG |
| Indicate a single nucleotide variant | chr1:5-5 T/A | chr1:4-5 T/A |

- 1-based : Single nucleotides, variant positions, or ranges are specified directly by their corresponding nucleotide numbers
  - GFF, SAM, VCF, Ensembl browser, ...
- 0-based: Single nucleotides, variant positions, or ranges are specified by the coordinates that flank them
  - BED, BAM, UCSC browser, ...

# Genome builds

## Reference Genome builds

Current human: GRCh38, hg38, b38  
alternates: GRCh38v2\_ccdg,  
GRCh38\_full\_analysis\_set\_plus\_decoy\_hla


Previous human: GRCh37, hg19, b37

Current mouse: GRCm38, mm10

For a detailed discussion of various human reference genome flavors refer here:

[https://pmbio.org/module-02-inputs/0002/02/01/Reference\\_Genome/](https://pmbio.org/module-02-inputs/0002/02/01/Reference_Genome/)

## Lift-over



# Variant shifting (alignment) and parsimony/trimming

Reference and alternative alleles of a CA short tandem repeat (STR)

REF  
ALT

GGGCACACACAGGG  
GGGCACACAGGG

← CA deletion from the reference

| Genome Reference  |                | Variant Call Format  |       |     |  |
|---|----------------|--|-------|-----|--|
|   | GGGCACACACAGGG | POS  | REF   | ALT |  |
| REF   | CA             | 8  | CA    | .   | Not left aligned and alternate allele is empty |
| ALT   | .              |  |       |     |  |
| REF   | CAC            | 6  | CAC   | C   | Not left aligned but parsimonious              |
| ALT   | C              |  |       |     |  |
| REF   | GCACA          | 3  | GCACA | GCA | Not right trimmed                              |
| ALT   | GCA            |  |       |     |  |
| REF   | GGCA           | 2  | GGCA  | GG  | Not left trimmed                               |
| ALT   | GG             |  |       |     |  |
| REF   | GCA            | 3  | GCA   | G   | Normalized (left aligned & parsimonious)       |
| ALT   | G              |  |       |     |  |
| Alleles represented against the human genome reference. Allele pairs are colored the same, all are representations of the same variant. |                | Alleles represented in Variant Call Format, all are representations of the same variant. |       |     |  |

**Parsimony:** representing variant in as few nucleotides as possible without reducing the length of any allele to 0

**Left (right) aligning =** shifting the start position of a variant as far to the left (right) as possible

# How should I sort my SAM/BAM file?

- Generally BAM files are sorted by position
  - This is for performance reasons
    - When sorted and indexed, arbitrary positions in a massive BAM file can be accessed rapidly
- Certain tools require a BAM sorted by read name
  - Usually this is when we need to easily identify both reads of a pair
    - The insert size between two reads may be large
    - In fusion detection we are interested in read pairs that map to different chromosomes



We are on a Coffee Break & Networking  
Session